



Using Rust in Radio Astronomy

Christopher Jordan

Curtin University, Australia



ASTRO 3D

**ARC CENTRE OF EXCELLENCE FOR
ALL SKY ASTROPHYSICS IN 3D**

**UNLOCKING THE UNIVERSE,
INSPIRING THE FUTURE**



ASTRO 3D

The MWA






ASTRO 3D

The MWA



ASTRO 3D


The MWA





<https://www.mwatelescope.org/>

The EoR

- Epoch of Reionisation (EoR)
 - First stars and galaxies
 - One of many MWA science projects
- Goal: Detect the EoR signature
 - Our approach: Neutral hydrogen
- Problems:
 - Extremely weak, distant signal
 - All other astronomical objects are “in the way”
 - Local radio interference
 - Earth’s ionosphere



The Data



antA	antB	ch0.xx	ch0.xy	ch0.yx	ch0.yy	ch1.xx	ch1.xy	ch1.yx	ch1.yy
0	0	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
0	1	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
0	2	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
0	3	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
1	1	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
1	2	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
1	3	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
2	2	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
2	3	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l
3	3	r,l	r,l	r,l	r,l	r,l	r,l	r,l	r,l

- All MWA data is stored in 3 dims:
 - time
 - baseline
 - channel
- A unit of data is 8 f32s
- A typical observation has:
 - 50 times (~ 2 minutes)
 - 8128 baselines
 - 768 channels
 - ~10 GB
- EoR needs ~ 1000 hours!

Rust solution: hyperdrive

- Previous MWA code was written in C, C++ & Python
 - All slow
 - All difficult to install
 - Despite supercomputers!
 - All onerous to run
 - Somewhat difficult to use
 - Poorly documented
- Rust rewrite: hyperdrive (search: mwa_hyperdrive)
 - Extremely easy to use
 - Easy to install
 - (Somewhat) well documented
 - Optionally uses a CUDA/HIP GPU
 - ... and very fast



https://github.com/MWATelescope/mwa_hyperdrive



Rust solution: hyperdrive

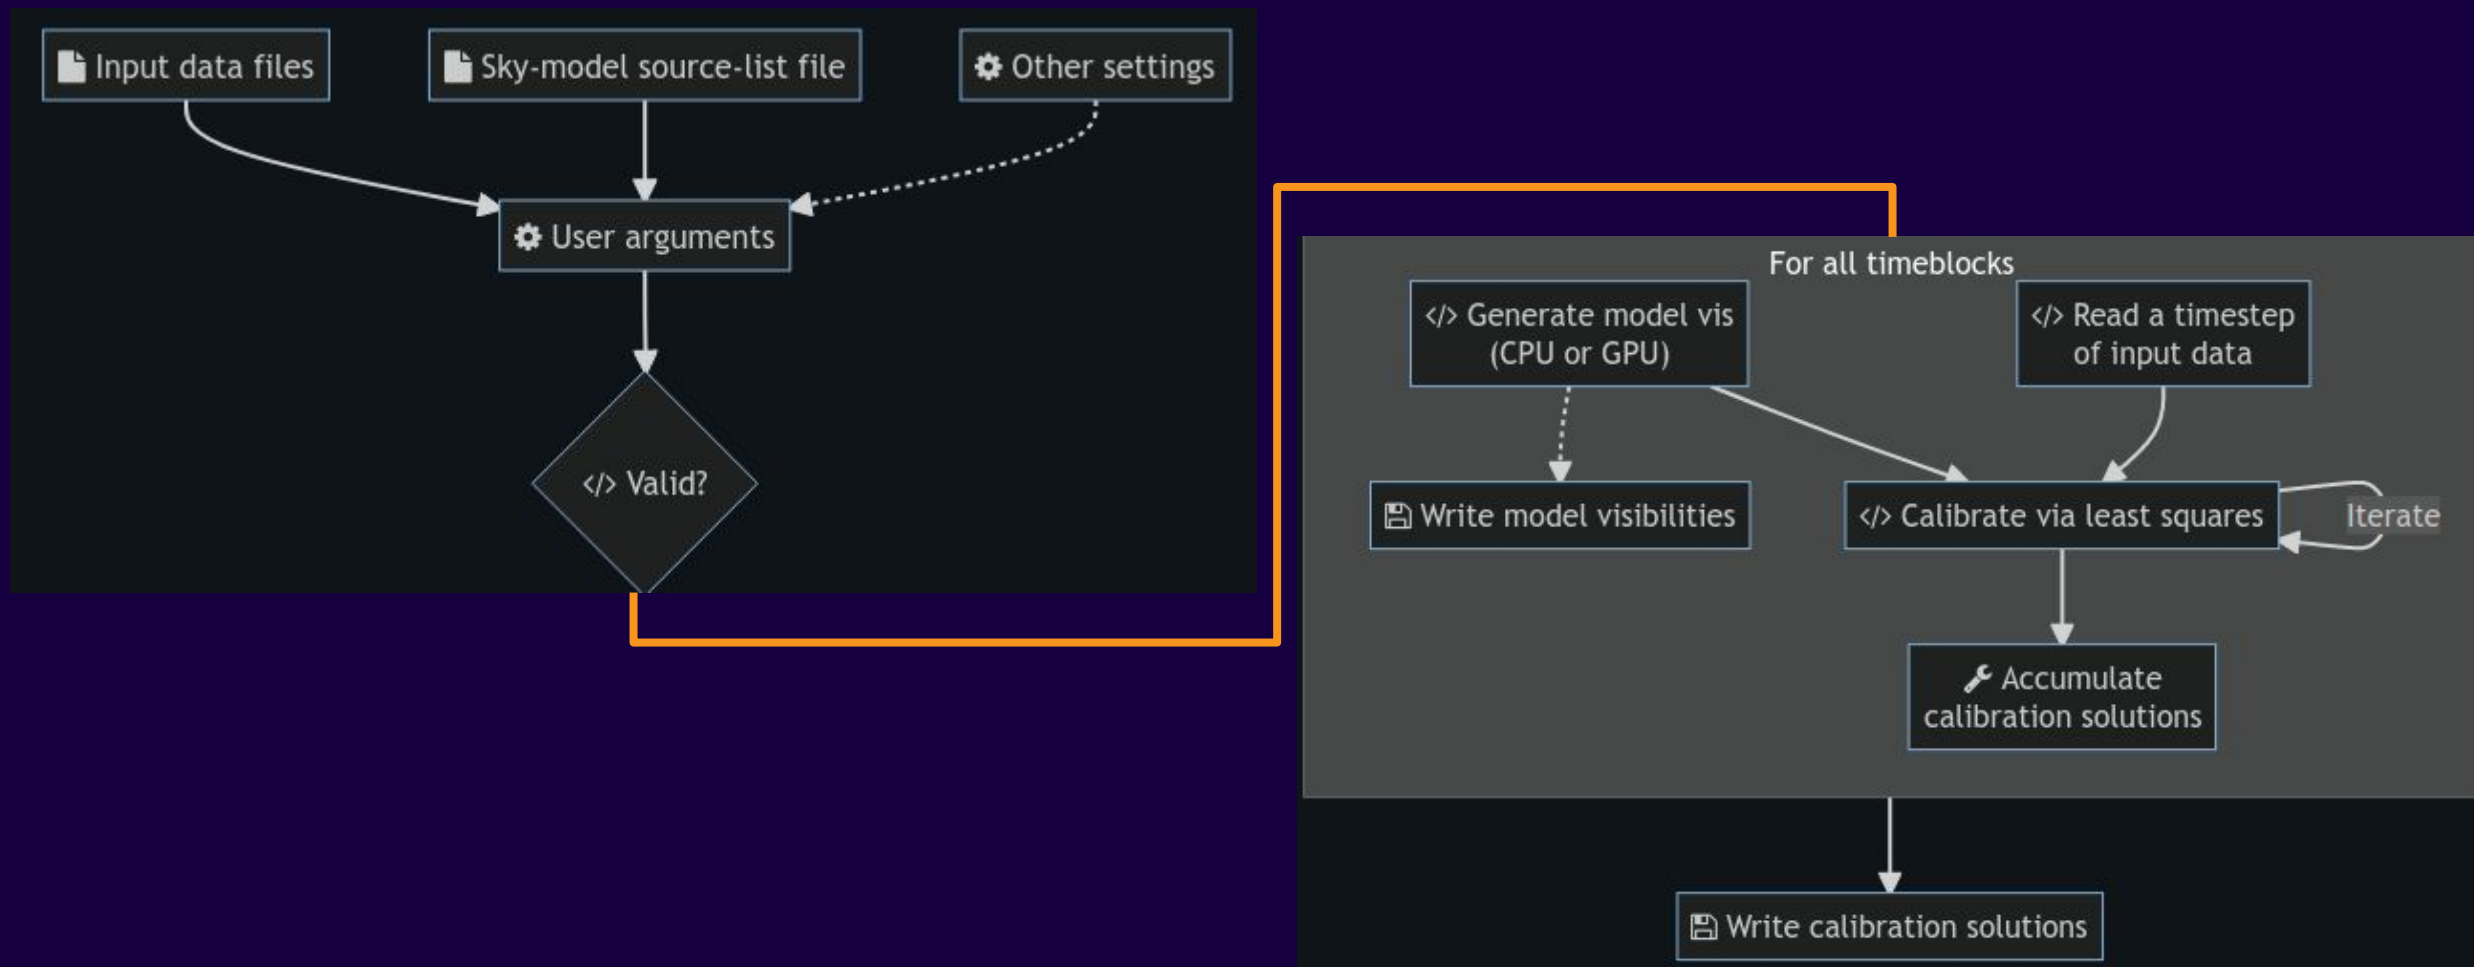
- Extremely useful crates:
 - ndarray (+ approx when testing)
 - hifitime
 - clap
 - bindgen (C dependencies + DIY GPU code)
 - rayon
 - crossbeam family
- Useful crates:
 - num-complex
 - thiserror
 - serde
 - indexmap
 - vec1
 - plotters
 - pyo3



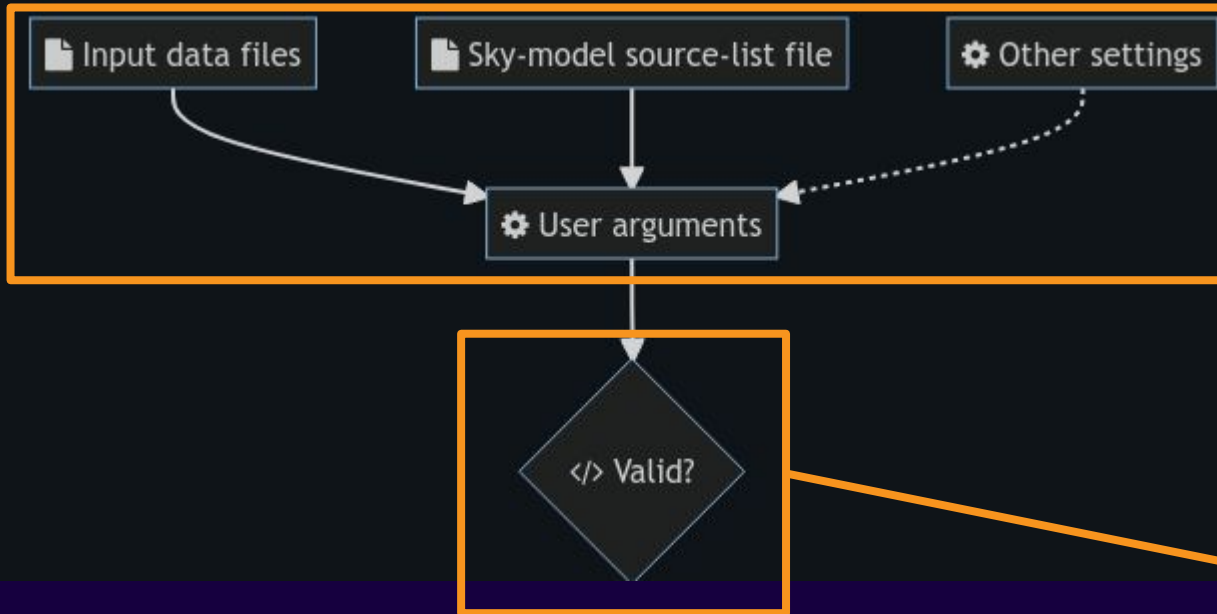
https://github.com/MWATelescope/mwa_hyperdrive



Programming approach



Programming approach



```
#[derive(Parser, Debug, Clone, Default, Serialize, Deserialize)]
struct DiCalCliArgs {
    #[clap(short='o', long="outputs", multiple_values(true), help = DI_SOLS_OUTPUT
    solutions: Option<Vec<PathBuf>>,

    /// The number of timesteps to average together during calibration. Also
    /// supports a target time resolution (e.g. 8s). If this is 0, then all data
    /// are averaged together. Default: 0. e.g. If this variable is 4, then we
    /// produce calibration solutions in timeblocks with up to 4 timesteps each.
    /// If the variable is instead 4s, then each timeblock contains up to 4s
    /// worth of data.
    #[clap(short, long, help_heading = "CALIBRATION")]
    timesteps_per_timeblock: Option<String>,
}
```

```
/// Parse the arguments into parameters ready for calibration.
fn parse(self) -> Result<DiCalParams, HyperdriveError> {
    debug!("{:#?}", self);

    let DiCalArgs {
        args_file: _,
        data_args,
        srclist_args,
        model_args,
        beam_args,
        calibration_args,
    } = self;

    let input_vis_params = data_args.parse("DI calibrating")?;
    let obs_context = input_vis_params.get_obs_context();
    let total_num_tiles = input_vis_params.get_total_num_tiles();
}
```

Programming approach

```
/// Use the [`DiCalParams`] to perform calibration and obtain solutions.
pub(crate) fn calibrate(&self) -> Result<CalibrationSolutions, super::DiCalibrateError> {
    // TODO: Fix.
    if self.freq_average_factor > 1 {
        panic!("Frequency averaging isn't working right now. Sorry!");
    }

    let CalVis {
        vis_data,
        vis_weights,
        vis_model,
    } = get_cal_vis(self, !self.no_progressBars?);
    assert_eq!(vis_weights.len_of(Axis(1)), self.baseline_weights.len());
```

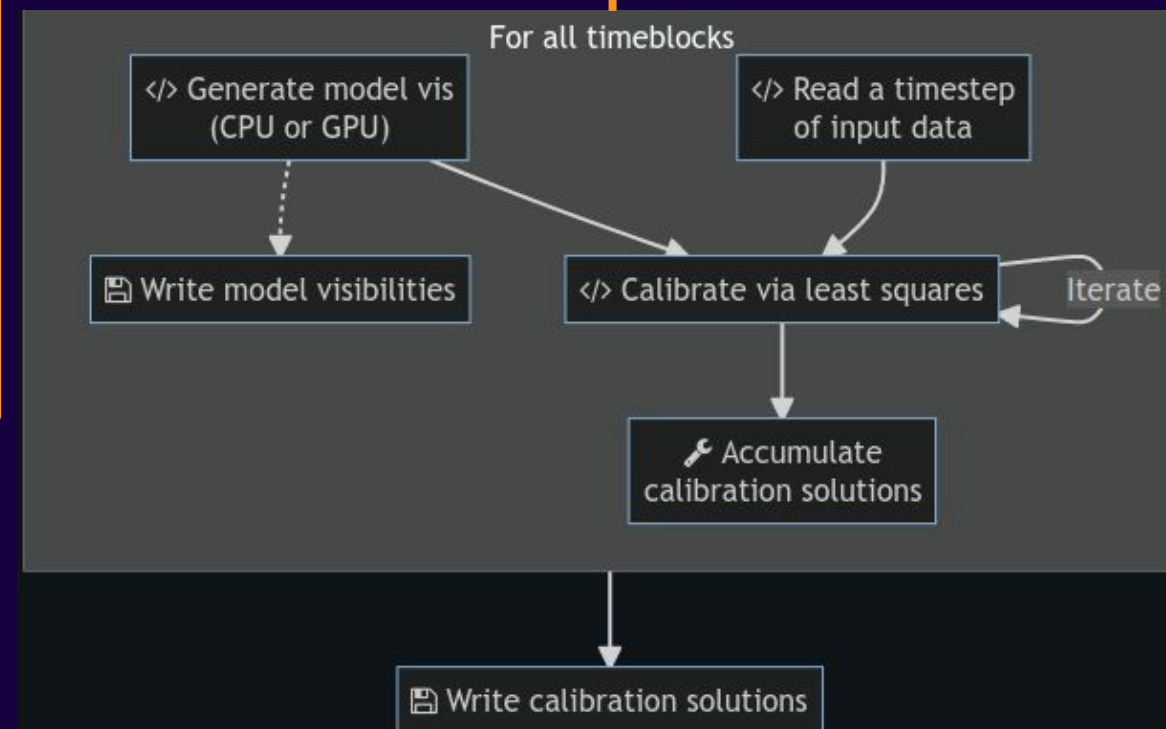
```
// Use a variable to track whether any threads have an issue.
let error = AtomicCell::new(false);
info!("Reading input data and sky modelling");
let scoped_threads_result = thread::scope(|scope| {
```

```
// Input visibility-data reading thread.
```

```
let data_handle = scope.spawn(|_| {
```

```
// Sky-model generation thread.
```

```
let model_handle = scope.spawn(|_| {
    defer_on_unwind! { error.store(true); }
```



Programming approach

- User command-line args are ingested
 - Everything is optional; let parsers determine what's needed
- Args are passed to parsers to create “params”
 - Everything is validated while parsed
 - Files readable/writable? etc.
 - Error messages are very detailed here
- Param structs are used to actually do work
 - Steps and dependencies are done on separate threads
 - Keeps memory low and disk IO/CPU high
 - Failure is less likely → performance is better
- Advantages
 - Modular
 - Correct
- Disadvantages
 - Test code can have lots of “ceremony” to get required types
 - (Probably) harsh on the compiler
 - “Interesting” code is “deep”; compile times mean it is slow to get feedback on new code

Programming approach: Use CPU or GPU

```
/// An object that simulates sky-model visibilities.
pub(crate) trait SkyModeller<'a> {
    /// Generate sky-model visibilities for a single timestep. The [UVW]
    /// coordinates used in generating the visibilities are returned.
    ///
    /// `vis_model_slice`: A mutable view into an `ndarray`. Rather than
    /// returning an array from this function, modelled visibilities are written
    /// into this array. This slice *must* have dimensions `[n1][n2]`, where
    /// `n1` is number of unflagged cross correlation baselines and `n2` is the
    /// number of unflagged frequencies.
    ///
    /// `timestamp`: The [hifitime::Epoch] struct used to determine what this
    /// timestep corresponds to.
    ///
    /// # Errors
    ///
    /// This function will return an error if there was a problem with
    /// beam-response calculation.
    fn model_timestep(
        &self,
        vis_model_slice: ArrayViewMut2<Jones<f32>>,
        timestamp: Epoch,
    ) -> Result<Vec<UVW>, BeamError>;
}
```

CPU

```
// Iterate over the unflagged baseline axis.
vis_model_slice
    .outer_iter_mut()
    .into_par_iter()
    .zip(uvws.par_iter())
    .enumerate()
    .for_each(|(i_baseline, (mut model_bl_axis, uvw_metres))| {
        let (i_tile1, i_tile2) = self.unflagged_baseline_to_tile_map[&i_baseline];

        // Unflagged fine-channel axis.
        model_bl_axis
            .iter_mut()
            .zip(fds.outer_iter())
            .zip(self.unflagged_fine_chan_freqs)
            .zip(beam_responses.slice(s![i_tile1, .., ..]).outer_iter())
            .zip(beam_responses.slice(s![i_tile2, .., ..]).outer_iter())
            .for_each(
                |((((model_vis, comp_fds), freq), tile1_beam), tile2_beam)| {
                    // Divide UVW by lambda to make UVW dimensionless.
                    let uvw = *uvw_metres * *freq / VEL_C;
                }
            );
    });
```


Programming approach: Use CPU or GPU

```
/// An object that simulates sky-model visibilities.
pub(crate) trait SkyModeller<'a> {
    /// Generate sky-model visibilities for a single timestep. The [UVW]
    /// coordinates used in generating the visibilities are returned.
    ///
    /// `vis_model_slice`: A mutable view into an `ndarray`. Rather than
    /// returning an array from this function, modelled visibilities are written
    /// into this array. This slice *must* have dimensions `[n1][n2]`, where
    /// `n1` is number of unflagged cross correlation baselines and `n2` is the
    /// number of unflagged frequencies.
    ///
    /// `timestamp`: The [hifitime::Epoch] struct used to determine what this
    /// timestep corresponds to.
    ///
    /// # Errors
    ///
    /// This function will return an error if there was a problem with
    /// beam-response calculation.
    fn model_timestep(
        &self,
        vis_model_slice: ArrayViewMut2<Jones<f32>>,
        timestamp: Epoch,
    ) -> Result<Vec<UVW>, BeamError>;
}
```


GPU

```
let cuda_status = cuda::model_points(
    &cuda::Points {
        num_power_law_points: self.point_power_law_radecs.len(),
        power_law_lmns: self.point_power_law_lmns.get_mut(),
        power_law_fds: self.point_power_law_fds.get_mut(),
        power_law_sis: self.point_power_law_sis.get_mut(),
    }
);
```

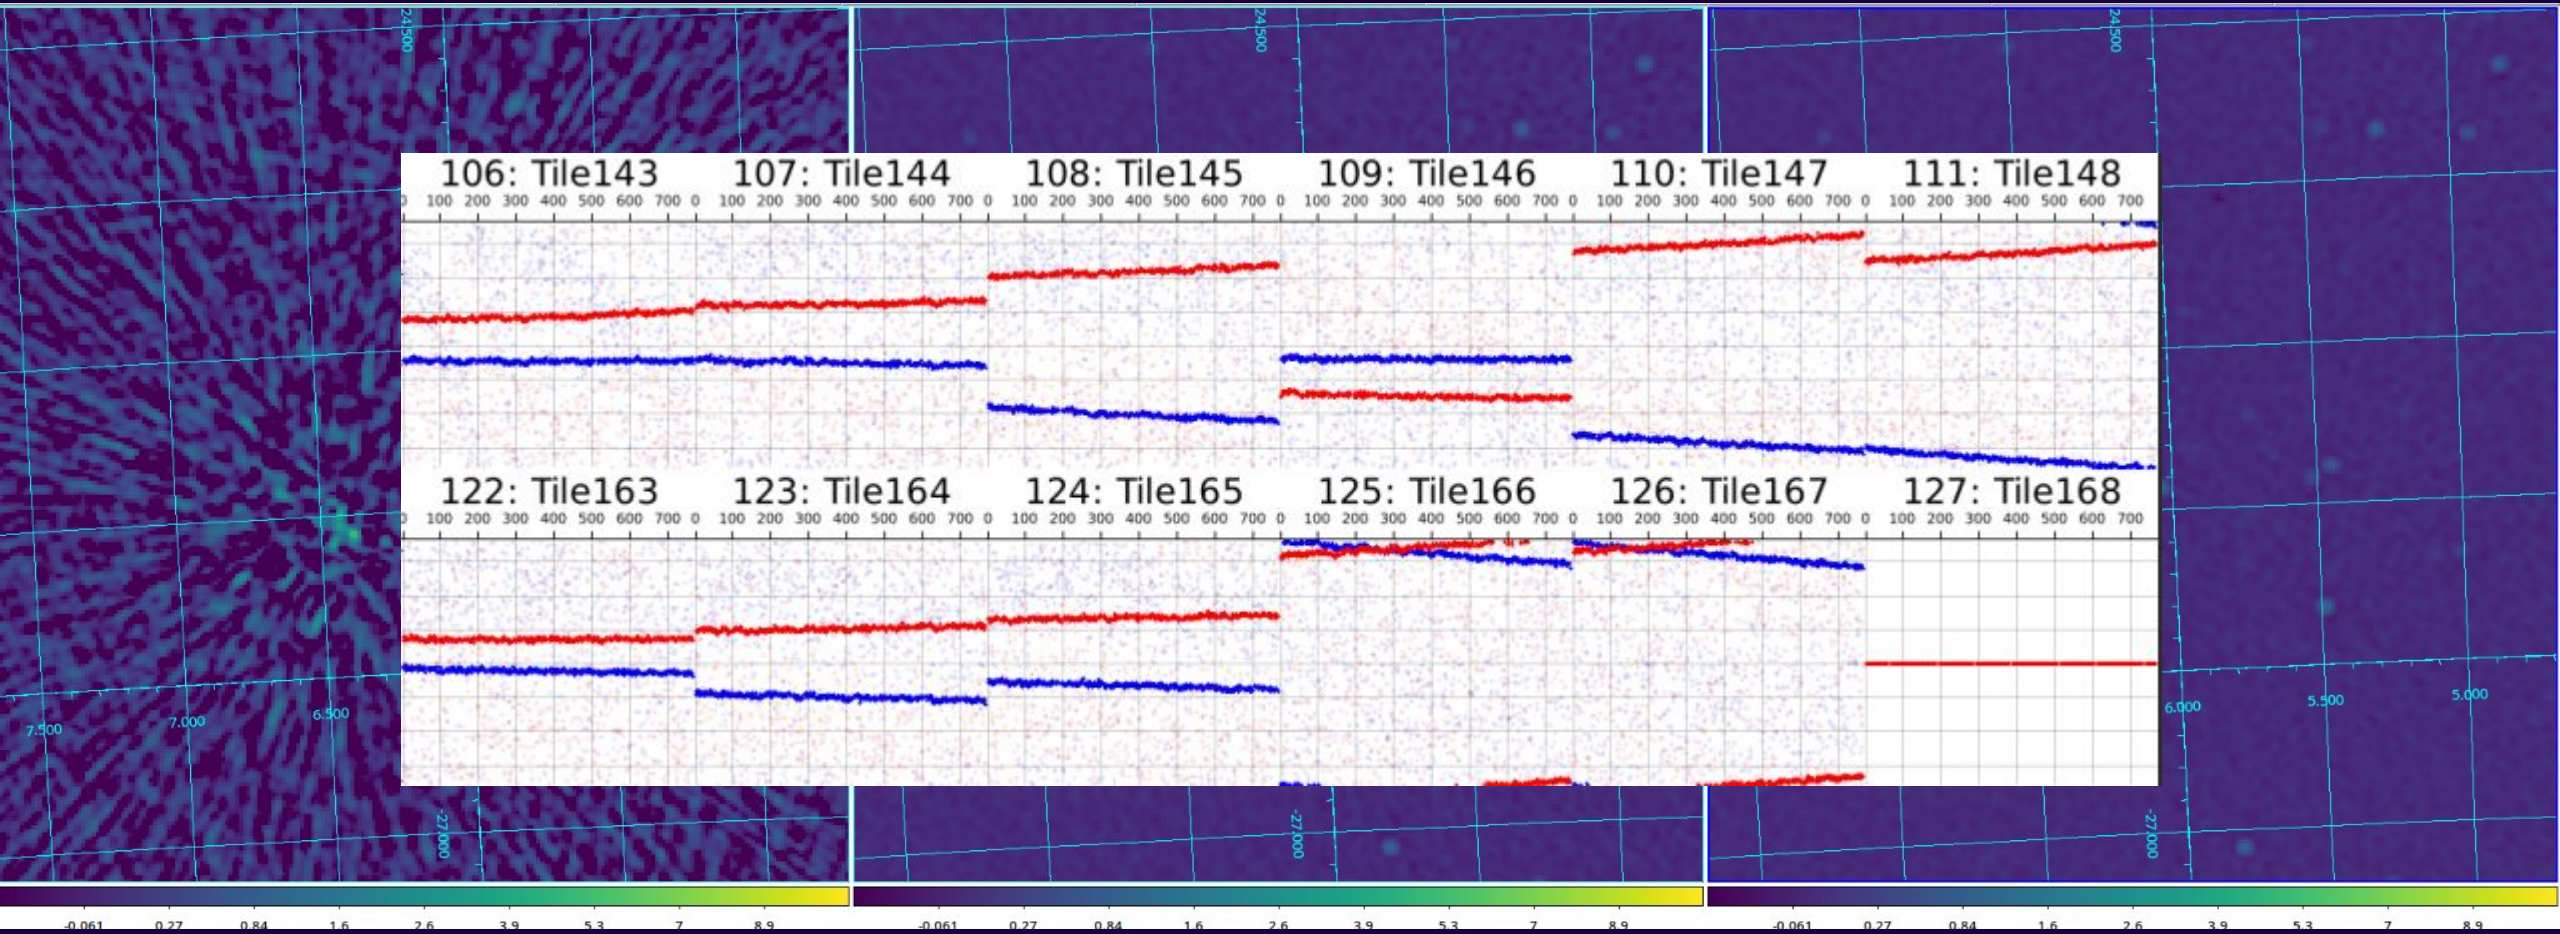
```
extern "C" int model_points(const Points *comps, const Addresses *a, const
    dim3 gridDim, blockDim;
    if (a->num_unique_beam_freqs == 0) {
        // Thread blocks are distributed by visibility (one visibility per
        // frequency and baseline).
        blockDim.x = 512;
        blockDim.y = 1;
        gridDim.x = (int)ceil((double)a->num_vis / (double)blockDim.x);
        gridDim.y = 1;

        model_points_kernel<<<gridDim, blockDim>>>>(a->num_freqs, a->num_vis,
            cudaCheck(cudaPeekAtLastError()));
    }
```



(Brief) MWA EoR results



(Brief) MWA EoR results



(Brief) MWA EoR results



```
$ hyperdrive di-calibrate -d data.uvfits data.metafits -s skymodel.json -o solutions.fits
```

```
$ hyperdrive plot solutions.fits
```


Using Rust has dramatically eased our:

- computational load,
- capacity to make changes quickly,
- maintenance burden, and
- ability to deploy/install on new machines.

But most importantly, Rust has enabled me to allow researchers to get on with their research.

I therefore strongly recommend using Rust for reliable, fast code.

Thanks!